

PROJECT(Merc_benz)

March 26, 2021

-importing libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import xgboost
from math import sqrt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
```

-Reading data

```
[2]: train_data= pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
[3]: train_data.head() #train data
```

```
[3]:
```

| | ID | y | X0 | X1 | X2 | X3 | X4 | X5 | X6 | X8 | ... | X375 | X376 | X377 | X378 | X379 | \ |
|---|----|--------|----|----|----|----|----|----|----|----|-----|------|------|------|------|------|---|
| 0 | 0 | 130.81 | k | v | at | a | d | u | j | o | ... | 0 | 0 | 1 | 0 | 0 | |
| 1 | 6 | 88.53 | k | t | av | e | d | y | l | o | ... | 1 | 0 | 0 | 0 | 0 | |
| 2 | 7 | 76.26 | az | w | n | c | d | x | j | x | ... | 0 | 0 | 0 | 0 | 0 | |
| 3 | 9 | 80.62 | az | t | n | f | d | x | l | e | ... | 0 | 0 | 0 | 0 | 0 | |
| 4 | 13 | 78.02 | az | v | n | f | d | h | d | n | ... | 0 | 0 | 0 | 0 | 0 | |

| | X380 | X382 | X383 | X384 | X385 |
|---|------|------|------|------|------|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 1 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 |

[5 rows x 378 columns]

```
[4]: target = train_data.y #assigning target variable
train_data = train_data.drop('y', axis=1) #features

print(train_data.shape)
```

```
print(target.shape)
```

```
(4209, 377)
(4209,)
```

```
[5]: target.isnull().sum() #checking missing values of target variable
```

```
[5]: 0
```

```
[6]: test_data.head() #test data
```

```
[6]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0    0    0    1    0    0
1    2   t  b  ai  a  d  b  g  y    0  ...    0    0    1    0    0    0
2    3  az  v  as  f  d  a  j  j    0  ...    0    0    0    1    0    0
3    4  az  l   n  f  d  z  l  n    0  ...    0    0    0    1    0    0
4    5   w  s  as  c  d  y  i  m    0  ...    1    0    0    0    0    0
```

```
      X382  X383  X384  X385
0         0         0         0         0
1         0         0         0         0
2         0         0         0         0
3         0         0         0         0
4         0         0         0         0
```

```
[5 rows x 377 columns]
```

```
[7]: train_data.shape
```

```
[7]: (4209, 377)
```

0.0.1 1. Removing the columns with variance is equal to zero

```
[8]: var_dt = train_data.var() #variance of columns
      var_dt
```

```
[8]: ID          5.941936e+06
      X10         1.313092e-02
      X11         0.000000e+00
      X12         6.945713e-02
      X13         5.462335e-02
      ...
      X380        8.014579e-03
      X382        7.546747e-03
      X383        1.660732e-03
      X384        4.750593e-04
```

```
X385      1.423823e-03
Length: 369, dtype: float64
```

```
[9]: var_dt[var_dt == 0].index #columns with variance equal to zero
```

```
[9]: Index(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293',
          'X297', 'X330', 'X347'],
          dtype='object')
```

```
[10]: train_data = train_data.drop(['ID', 'X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347'], axis=1)
train_data.shape
```

[10]: (4209, 364)

```
[11]: test_data = test_data.drop(['ID', 'X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347'], axis=1)
test_data.shape
```

```
[11]: (4209, 364)
```

0.0.2 2. Checking for null and unique values for test and train sets.

```
[12]: print(list(train_data.isnull().sum())) #missing values of train data
```

[illegible]

```
[13]: print(list(test_data.isnull().sum())) #missing values of test data
```

[illegible]

[illegible]

```
[14]: print(len(train_data.X0.unique())) #no of unique values for every column of
      ↪ train_data before analysis
      print(len(train_data.X1.unique()))
      print(len(train_data.X2.unique()))
      print(len(train_data.X3.unique()))
      print(len(train_data.X4.unique()))
      print(len(train_data.X5.unique()))
      print(len(train_data.X6.unique()))
      print(len(train_data.X8.unique()))
```

47
27
44
7
4
29
12
25

```
[15]: print(len(test_data.X0.unique())) #no of unique values for every column of
      ↪ test_data before analysis
      print(len(test_data.X1.unique()))
      print(len(test_data.X2.unique()))
      print(len(test_data.X3.unique()))
      print(len(test_data.X4.unique()))
      print(len(test_data.X5.unique()))
      print(len(test_data.X6.unique()))
      print(len(test_data.X8.unique()))
```

49
27
45
7
4
32
12
25

```
[16]: # checking different values for same column X0 in train and test data.
set_a=set(test_data.X0.drop_duplicates())
set_b=set(train_data.X0.drop_duplicates())
print(set_b-set_a) #values in train but not in test
print(set_a-set_b) #values in test but not in train
```

```
{'ab', 'q', 'aa', 'ac'}
{'av', 'ae', 'bb', 'p', 'ag', 'an'}
```

```
[17]: set_a == set_b
```

```
[17]: False
```

```
[18]: test_data['X0'].replace({"av":"q","p":"ac","ae":"aa","ag":"ab"},inplace=True)
```

```
[19]: test_data.drop(test_data.index[test_data['X0'] == 'an'], inplace=True)
test_data.drop(test_data.index[test_data['X0'] == 'bb'], inplace=True)
set_a == set_b
```

```
[19]: False
```

```
[20]: # checking different values for same column X1 in train and test data.
set_a=set(test_data.X1.drop_duplicates())
set_b=set(train_data.X1.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a ==set_b
```

```
set()
set()
```

```
[20]: True
```

```
[21]: # checking different values for same column X2 in train and test data.
set_a=set(test_data.X2.drop_duplicates())
set_b=set(train_data.X2.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a == set_b
```

```
{'l', 'ar', 'av', 'c', 'o', 'aa'}
{'ab', 'ad', 'u', 'ax', 'w', 'aj'}
```

```
[21]: False
```

```
[22]: test_data['X2'].replace({"u":"aa","w":"l","ab":"o","aj":"c","ad":
↪ "ar"},inplace=True)
train_data['X2'].replace({"av":"aa"},inplace=True)
```

```
[23]: test_data.drop(test_data.index[test_data['X2'] == 'ax'], inplace=True)
```

```
[24]: # checking different values for same column X3 in train and test data.
set_a=set(test_data.X3.drop_duplicates())
set_b=set(train_data.X3.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a ==set_b
```

```
set()
set()
```

```
[24]: True
```

```
[25]: # checking different values for same column X4 in train and test data.
set_a=set(test_data.X4.drop_duplicates())
set_b=set(train_data.X4.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a ==set_b
```

```
set()
set()
```

```
[25]: True
```

```
[26]: # checking different values for same column X5 in train and test data.
set_a=set(test_data.X5.drop_duplicates())
set_b=set(train_data.X5.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a ==set_b
```

```
{'u'}
{'a', 'z', 't', 'b'}
```

```
[26]: False
```

```
[27]: test_data['X5'].replace({"b":"u"},inplace=True)
```

```
[28]: test_data.drop(test_data.index[test_data['X5'] == 'a'], inplace=True)
test_data.drop(test_data.index[test_data['X5'] == 't'], inplace=True)
test_data.drop(test_data.index[test_data['X5'] == 'z'], inplace=True)
```

```
[29]: # checking different values for same column X6 in train and test data.
set_a=set(test_data.X6.drop_duplicates())
set_b=set(train_data.X6.drop_duplicates())
print(set_b-set_a)#values in train but not in test
```

```
print(set_a-set_b)#values in test but not in train
set_a == set_b
```

```
set()
```

```
set()
```

[29]: True

```
[30]: # checking different values for same column X8 in train and test data.
set_a=set(test_data.X8.drop_duplicates())
set_b=set(train_data.X8.drop_duplicates())
print(set_b-set_a)#values in train but not in test
print(set_a-set_b)#values in test but not in train
set_a ==set_b
```

```
set()
```

```
set()
```

[30]: True

```
[31]: # train data
cat_data = train_data.select_dtypes('object')
int_data = train_data.select_dtypes('int64')

# test data
cat_test = test_data.select_dtypes('object')
int_test = test_data.select_dtypes('int64')
```

```
[32]: print(len(cat_data.X0.unique())) #no of unique values for every column of
      ↪train_data after analysis
print(len(cat_data.X1.unique()))
print(len(cat_data.X2.unique()))
print(len(cat_data.X3.unique()))
print(len(cat_data.X4.unique()))
print(len(cat_data.X5.unique()))
print(len(cat_data.X6.unique()))
print(len(cat_data.X8.unique()))
```

```
47
```

```
27
```

```
43
```

```
7
```

```
4
```

```
29
```

```
12
```

```
25
```

```
[33]: print(len(cat_test.X0.unique())) #no of unique values for every column of
      ↪ test_data after analysis
print(len(cat_test.X1.unique()))
print(len(cat_test.X2.unique()))
print(len(cat_test.X3.unique()))
print(len(cat_test.X4.unique()))
print(len(cat_test.X5.unique()))
print(len(cat_test.X6.unique()))
print(len(cat_test.X8.unique()))
```

```
47
27
43
7
4
29
12
25
```

```
[34]: print(train_data.shape)
      print(test_data.shape)
```

```
(4209, 364)
(4203, 364)
```

0.0.3 3. Encoding Categorical Variables.

```
[95]: from sklearn.preprocessing import LabelEncoder
      le = LabelEncoder()
      train_data['X0'] = le.fit_transform(train_data['X0'])
      test_data['X0'] = le.transform(test_data['X0'])
```

```
[96]: train_data['X1'] = le.fit_transform(train_data['X1'])
      test_data['X1'] = le.transform(test_data['X1'])
```

```
[97]: train_data['X2'] = le.fit_transform(train_data['X2'])
      test_data['X2'] = le.transform(test_data['X2'])
```

```
[98]: train_data['X3'] = le.fit_transform(train_data['X3'])
      test_data['X3'] = le.transform(test_data['X3'])
```

```
[99]: train_data['X4'] = le.fit_transform(train_data['X4'])
      test_data['X4'] = le.transform(test_data['X4'])
```

```
[100]: train_data['X5'] = le.fit_transform(train_data['X5'])
       test_data['X5'] = le.transform(test_data['X5'])
```



```
[101]: train_data['X6'] = le.fit_transform(train_data['X6'])
test_data['X6'] = le.transform(test_data['X6'])
```

```
[102]: train_data['X8'] = le.fit_transform(train_data['X8'])
test_data['X8'] = le.transform(test_data['X8'])
```

```
[103]: train_data.head()
```

```
[103]:
```

| | X0 | X1 | X2 | X3 | X4 | X5 | X6 | X8 | X10 | X12 | ... | X375 | X376 | X377 | X378 | \ |
|---|----|----|----|----|----|----|----|----|-----|-----|-----|------|------|------|------|---|
| 0 | 32 | 23 | 17 | 0 | 3 | 24 | 9 | 14 | 0 | 0 | ... | 0 | 0 | 1 | 0 | |
| 1 | 32 | 21 | 1 | 4 | 3 | 28 | 11 | 14 | 0 | 0 | ... | 1 | 0 | 0 | 0 | |
| 2 | 20 | 24 | 33 | 2 | 3 | 27 | 9 | 23 | 0 | 0 | ... | 0 | 0 | 0 | 0 | |
| 3 | 20 | 21 | 33 | 5 | 3 | 27 | 11 | 4 | 0 | 0 | ... | 0 | 0 | 0 | 0 | |
| 4 | 20 | 23 | 33 | 5 | 3 | 12 | 3 | 13 | 0 | 0 | ... | 0 | 0 | 0 | 0 | |

| | X379 | X380 | X382 | X383 | X384 | X385 |
|---|------|------|------|------|------|------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 1 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 |

[5 rows x 364 columns]

```
[104]: test_data.head()
```

```
[104]:
```

| | X0 | X1 | X2 | X3 | X4 | X5 | X6 | X8 | X10 | X12 | ... | X375 | X376 | X377 | X378 | \ |
|---|----|----|----|----|----|----|----|----|-----|-----|-----|------|------|------|------|---|
| 1 | 40 | 3 | 7 | 0 | 3 | 24 | 6 | 24 | 0 | 0 | ... | 0 | 0 | 1 | 0 | |
| 4 | 43 | 20 | 16 | 2 | 3 | 28 | 8 | 12 | 0 | 0 | ... | 1 | 0 | 0 | 0 | |
| 5 | 45 | 1 | 7 | 4 | 3 | 27 | 6 | 18 | 0 | 0 | ... | 1 | 0 | 0 | 0 | |
| 6 | 44 | 3 | 3 | 3 | 3 | 27 | 3 | 24 | 0 | 0 | ... | 0 | 0 | 0 | 0 | |
| 7 | 27 | 20 | 3 | 2 | 3 | 12 | 3 | 0 | 0 | 0 | ... | 0 | 0 | 1 | 0 | |

| | X379 | X380 | X382 | X383 | X384 | X385 |
|---|------|------|------|------|------|------|
| 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 |
| 5 | 0 | 0 | 0 | 0 | 0 | 0 |
| 6 | 0 | 1 | 0 | 0 | 0 | 0 |
| 7 | 0 | 0 | 0 | 0 | 0 | 0 |

[5 rows x 364 columns]

```
[105]: #splitting the train_data into test and train to testing the accuracy..
x_train,x_test,y_train,y_test = train_test_split(train_data,target, test_size=0.
↪25, random_state=50)
```

```
[106]: print(x_train.shape)
       print(x_test.shape)
       print(y_train.shape)
       print(y_test.shape)
```

```
(3156, 364)
(1053, 364)
(3156,)
(1053,)
```

```
[107]: from sklearn.preprocessing import StandardScaler
       from sklearn.decomposition import PCA as sklearnPCA
       import matplotlib.pyplot as plt
```

```
[108]: sc = StandardScaler() #scaling data
```

```
[109]: sc.fit(x_train)
```

```
[109]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
[110]: X_std_train = sc.transform(x_train)
       X_std_test = sc.transform(x_test)
```

```
[111]: X_test = sc.transform(test_data) #scaling test_data
```

0.0.4 4. Performing Dimensionality Reduction.

```
[112]: pca = sklearnPCA() # dimensionality reduction with PCA
```


```
[152]: sklearn_pca = pca.fit(X_std_train)
       train_pca = pd.DataFrame(sklearn_pca.transform(X_std_train))
       test_pca = pd.DataFrame(sklearn_pca.transform(X_std_test))
```

```
[159]: X_test_pca = pd.DataFrame(sklearn_pca.transform(X_test)) #dimensionality
       ↪reduction on test_data
```

```
[154]: var_per = sklearn_pca.explained_variance_ratio_ #variance per component
       cum_var_per = sklearn_pca.explained_variance_ratio_.cumsum() #cummulative
       ↪variance
```

```
[155]: plt.figure(figsize=(30,10))
       ind = np.arange(len(var_per))
       plt.bar(ind,var_per)
       plt.xlabel('n_components')
       plt.ylabel('variance')
```

```
[155]: Text(0, 0.5, 'variance')
```



```
[156]: len(cum_var_per[cum_var_per < 0.95]) # no of components explaining 95% variance.
```

```
[156]: 140
```

```
[157]: n_comp=len(cum_var_per[cum_var_per <= 0.95])
print("Keeping 95% Info with ",n_comp," components")
sklearn_pca = sklearnPCA(n_components=n_comp)
train_pca = pd.DataFrame(sklearn_pca.fit_transform(X_std_train))
test_pca = pd.DataFrame(sklearn_pca.transform(X_std_test))
print("Shape before PCA for Train: ",X_std_train.shape)
print("Shape after PCA for Train: ",train_pca.shape)
print("Shape before PCA for Test: ",X_std_test.shape)
print("Shape after PCA for Test: ",test_pca.shape)
```

```
Keeping 95% Info with 140 components
Shape before PCA for Train: (3156, 364)
Shape after PCA for Train: (3156, 140)
Shape before PCA for Test: (1053, 364)
Shape after PCA for Test: (1053, 140)
```

```
[160]: X_test_pca.shape #shape after pca on test_data
```

```
[160]: (4203, 140)
```

-Fitting model

```
[161]: lin = LinearRegression()
```

```
[162]: lin.fit(train_pca,y_train)
```

```
[162]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[163]: y_pred = lin.predict(test_pca)
```

```
[164]: print(sqrt(mean_squared_error(y_test, y_pred)))
```

8.347089870204956

```
[165]: print(r2_score(y_test, y_pred))
```

0.5453912425394598

```
[166]: lin.score(test_pca, y_test)
```

[166]: 0.5453912425394598

```
[167]: lin.score(train_pca, y_train)
```

[167]: 0.5635937366469262

```
[183]: # Let's try XGboost algorithm to see if we can get better results
xgb = xgboost.XGBRegressor(objective='reg:linear', learning_rate=0.08)
```

```
[184]: xgb.fit(train_pca,y_train)
```

[16:36:25] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear is now deprecated in favor of reg:squarederror.

```
[184]: XGBRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
importance_type='gain', interaction_constraints=None,
learning_rate=0.08, max_delta_step=0, max_depth=6,
min_child_weight=1, missing=nan, monotone_constraints=None,
n_estimators=100, n_jobs=0, num_parallel_tree=1,
objective='reg:linear', random_state=0, reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, subsample=1, tree_method=None,
validate_parameters=False, verbosity=None)
```

```
[185]: predictions = xgb.predict(test_pca)
xgb.score(test_pca, y_test)
```

[185]: 0.5263172696229608

```
[186]: xgb.score(train_pca,y_train)
```

[186]: 0.8785191593221224

```
[187]: print(r2_score(y_test, predictions))
```

0.5263172696229608

```
[188]: xgrf = xgboost.XGBRFRegressor()
```

```
[189]: xgrf.fit(train_pca,y_train)
```

```
[189]: XGBRFRegressor(base_score=0.5, booster=None, colsample_bylevel=1,  
                    colsample_bynode=0.8, colsample_bytree=1, gamma=0, gpu_id=-1,  
                    importance_type='gain', interaction_constraints=None,  
                    learning_rate=1, max_delta_step=0, max_depth=6,  
                    min_child_weight=1, missing=nan, monotone_constraints=None,  
                    n_estimators=100, n_jobs=0, num_parallel_tree=100,  
                    objective='reg:squarederror', random_state=0, reg_alpha=0,  
                    reg_lambda=1e-05, scale_pos_weight=1, subsample=0.8,  
                    tree_method=None, validate_parameters=False, verbosity=None)
```

```
[190]: predictions = xgrf.predict(test_pca)  
xgrf.score(test_pca, y_test)
```

```
[190]: 0.4960210454925235
```

```
[179]: xgrf.score(train_pca, y_train)
```

```
[179]: 0.6129669157901589
```

```
[ ]:
```

```
[ ]:
```